

OpenEd Trust, Safety, and Content QA Document

Purpose

Trust is central to OpenEd. A free learning platform can degrade quickly if it accepts low-quality courses, unsafe content, broken sources, generic AI-generated lessons, or misleading proof.

Trust pillars

1. Source transparency.
2. Course quality review.
3. Educator accountability.
4. AI tutor safety.
5. Assessment integrity.
6. Learner privacy.
7. Moderation and reporting.

Course QA rubric

Content completeness

- Course has clear audience.
- Course has prerequisites.
- Course has learning outcomes.
- Every lesson has objective.
- Every lesson has summary.
- Every lesson has practice or knowledge check.
- Every module has artifact or assessment.

Source quality

- Claims have sources where needed.
- Sources are relevant.

- Sources are not fabricated.
- Sources are marked required/optional.
- Advanced resources are separated.
- Broken links are flagged.

Assessment quality

- Quizzes have explanations.
- Wrong answers have feedback.
- Rubrics have criteria and descriptors.
- Artifacts have examples/non-examples.
- Proof claims are not overstated.

AI tutor readiness

- Lesson has tutor context.
- Suggested prompts exist.
- Restricted topics are flagged.
- Tutor should know when to refuse/escalate.

Accessibility

- Video has transcript/summary.
- Diagrams have text summaries.
- Contrast passes.
- Mobile layout works.
- Language is not unnecessarily complex.

Course publishing states

```
draft -> qa_ready -> submitted -> in_review -> changes_requested -> approved -> published  
-> monitored -> archived
```

Moderation policy

Report types:

- Incorrect course content.
- Broken source.

- Harmful or unsafe content.
- Tutor hallucination.
- Cheating request.
- Harassment/abuse.
- IP/copyright concern.
- Privacy concern.

Severity:

Severity	Response
Low	Batch review.
Medium	Review within 72 hours.
High	Review within 24 hours.
Critical	Immediate restriction or temporary removal.

Safety-sensitive content

OpenEd can host serious courses on technical or sensitive topics, but must use safe teaching methods.

For high-risk domains:

- Use abstracted/sandboxed examples.
- Avoid operational harmful steps.
- Use defensive framing.
- Add source and safety review.
- Restrict tutor behavior.
- Escalate to OpenEd Team review.

AI tutor safety controls

- Source-grounded mode.
- Refusal for unsafe requests.
- No direct answer during locked assessment.
- Report answer button.
- Tutor failure log.
- Prompt injections handled as untrusted user content.

Privacy controls

- No BYOK key storage on server.
- Optional local browser storage with warning.
- No sensitive learner data in tutor logs by default.
- Private artifacts remain private unless shared.
- Admin access is audited.

Content QA automation

Automated checks can flag:

- Missing objective.
- Missing source.
- Broken link.
- Duplicate lesson text.
- No quiz feedback.
- No rubric.
- Missing transcript.
- Unsupported claim language.
- Unsafe terms requiring human review.

Automation should not be final authority. It supports OpenEd Team review.

Quality badge model

Course badges:

- Source-mapped.
- Tutor-ready.
- Artifact-based.
- Rubric-reviewed.
- Accessibility-checked.
- OpenEd reviewed.

Do not create misleading badges such as "job guaranteed" or "expert certified" without evidence.

Trust failure modes

Failure	Mitigation
Fake sources	Source verification and reviewer checks.
Generic AI lessons	Duplicate/generic text detection and quality review.
Harmful course content	Restricted-topic policies and human review.
Tutor hallucination	Source grounding and report workflow.
Misleading certificate	Proof ladder and transparent evidence.
Educator abuse	Reporting, suspension, audit logs.

Source notes

This document is grounded in the current OpenEd/FME project context and the following reference set where applicable:

- Uploaded Frontier Model Evaluations prototype code bundle, latest inspected bundle: Design System for Frontier Model Evaluation (5).zip.
- Uploaded Frontier Model Evaluations 51-hour curriculum and Module 1 foundations files.
- Coursera, Udemy, DataCamp, Moodle, Khanmigo, Duolingo public documentation and product pages.
- Supabase Auth and Row Level Security documentation.
- OpenAI API key safety guidance, MDN Web APIs, and YouTube IFrame Player API documentation.
- Bloom's 2 sigma problem, intelligent tutoring systems research, CourseAssist, LearnLM-Tutor, DCCI, and UNESCO guidance for generative AI in education.

Document version: OpenEd v1 planning pack, 2026-06-17.